

# Material Summary: Working with Images

## 1. Image Processing

### 1.1 Loading and Inspecting Images

- There are many ways to read an image
- One of the easiest is using scikit-image

```
from skimage.io import imread  
tiger_image = imread("tiger.jpg")
```

- Displaying the image


```
plt.imshow(tiger_image)
```

- The image is actually a **matrix of pixels**
  - Each pixel is an array of three values: **R, G, B**  $\in [0; 255]$
  - Grayscale images only have **one value per pixel**
- Most image **processing algorithms** are easier to understand on grayscale images

```
red = tiger_image[:, :, 0]  
green = tiger_image[:, :, 1]  
blue = tiger_image[:, :, 2]
```

### 1.2 Image Histogram

- As usual, histograms tell us how the values are distributed
- How many dark values, how many light values
- Maximum brightness, peaks, etc.



- Histograms need to have a **single variable**
  - Take each channel separately, e.g., red
  - Convert the **2D matrix to 1D array**: `image.ravel()`
  - Show the histogram as usual
    - It's common to use 256 bins

```
plt.hist(red.ravel(), bins = 256, color = "red")  
plt.show()
```

- We can also plot **all channels on a single histogram**

### 1.3 Converting to Grayscale

- Sometimes working **per channel is not necessary**
- We can combine **all three channels** and get a grayscale image
- Simplest way: **get the mean of all values**

```
tiger_grayscale = np.mean(tiger_image, axis = 2)
```


- Better way: **use coefficients for each channel**
  - The human eye **discerns colors differently**
  - Were more sensitive to green colors
  - Some formulas are given [here](#)

```
tiger_grayscale = 0.299 * red + 0.587 * green + 0.114 * blue
```

- Depending on the image, the differences **may or may not be easy to see**
- For art purposes, we can experiment with our own **coefficients for combining all channels**


### 1.4 Convolution

- **Convolution kernel (filter)**
  - A small, usually 3x3, matrix of numbers
- Convolution process
  - **Input:** image, kernel
  - **Output:** new image



- Combining the image and a kernel:
  - Apply the kernel **over each pixel**
  - Multiply the values **element-wise** (Hadamard product)
  - Sum **all values**
  - Assign the **sum to the corresponding pixel** in the output image
    - Image corners are treated in different ways, not really important how
- The choice of kernel depends what the **output image will represent**


```
from scipy.ndimage.filters import convolve  
convolve(image, kernel)
```



- Example: **box blur**


```
box_blur_kernel = np.array([
    [1, 1, 1],
    [1, 1, 1],
    [1, 1, 1]
]) / 9

blurred = convolve(tiger_grayscale, box_blur_kernel)
plt.imshow(tiger_grayscale[150:250, 300:400], cmap = "gray")
plt.show()
plt.imshow(blurred[150:250, 300:400], cmap = "gray")
plt.show()
```




## 1.5 Image Morphology

- Four main operations (see [this](#) tutorial)
  - **Dilation**
  - **Erosion**
  - **Opening**
  - **Closing**
- A simple series of **algorithms for image transformation**
- Basic methodology
  - Choose a structuring element (e.g., 2x2 square or cross)
  - Move the element around the image
  - Apply an operation
- **Input: binary image**
  - Pixel values **0** and **1**, not **[0; 255]**
  - This is called **thresholding**
- **Output: transformed image**






First get **all values** inside the structuring element

- **Erosion:** replace all values with the min value
  - Strips away a layer of pixels
  - Holes become larger
  - Small regions are eliminated
- **Dilation:** replace all values with the max value



- Adds a layer of pixels
- Gaps become smaller
- Small gaps are filled in
- **Opening: erosion followed by dilation**
  - Pixels which survived erosion are restored to their original size
  - Opens up a gap between two objects connected by thin bridges
- **Closing: dilation followed by erosion**
  - Fills in holes in the regions while keeping the initial region sizes



## 1.6 Other Operations on Images

- Matrix operations – **pixel-wise**
  - One image:
    - **Addition, Gain, Negative**
    - **Resampling, Cutting**
  - Transformations – **perspective, warp, etc.**
  - Two (or more) images:
    - **Addition (multiple exposure)**
    - **Subtraction (difference)**
    - **Division (normalization)**
    - **Averaging**
- Thresholding (**usually 2 levels**)
- Fourier **transform, filtering and convolution**
- Contrast **enhancement**
- Histogram **equalization**
- Stacking (many **2D images** ⇒ **one 3D image**)
- Analysis:
  - **Measurements, Segmentation, Object extraction / Identification**
  - **Enhancements, Inpainting**